

Here is a screenshot showing both hi and sup being recognized with high confidence.

```
hi: 0.857476  
sup: 0.142524
```

```
hi: 0.552441  
sup: 0.447559
```

```
hi: 0.263683  
sup: 0.736317
```

Software issue

The model being trained was a classification model, but it used MSE loss and MAE metrics.

This was easy to change, I just changed the loss and metrics in `model.compile()` to be categorical cross entropy and accuracy. I also had to change the metric to plot into accuracy so that the plots say “accuracy” in the title and the y-axis label.